

1-classifier-comparison

November 22, 2018

0.1 Classifier Comparison

In this notebook, we compare some of traditional machine learning algorithms on movie review datasets.

```
In [1]: import sys, argparse
        sys.path.append('.')
        import helper
        import numpy as np
        from sklearn.feature_extraction.text import TfidfVectorizer
        from sklearn.neighbors import KNeighborsClassifier
        from sklearn.svm import LinearSVC
        from sklearn.linear_model import SGDClassifier, LogisticRegression
        from sklearn.neural_network import MLPClassifier
        from sklearn.tree import DecisionTreeClassifier
        from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
        from sklearn.naive_bayes import GaussianNB
        from sklearn import metrics
```

0.1.1 1. Data preprocessing

```
In [2]: pos_examples = [s.decode("utf-8", "ignore").strip() for s in list(open(helper.mr_pos_data, 'r'))]
        neg_examples = [s.decode("utf-8", "ignore").strip() for s in list(open(helper.mr_neg_data, 'r'))]
        pos_nums, neg_nums = len(pos_examples), len(neg_examples)
        x = pos_examples + neg_examples
        x = [helper.clean_str(sentence) for sentence in x]
        pos_labels = [1 for _ in range(pos_nums)]
        neg_labels = [0 for _ in range(neg_nums)]
        y = pos_labels + neg_labels
        x, y = np.array(x), np.array(y)
        x_train, y_train, x_dev, y_dev = helper.split_train_dev(x, y)

In [3]: tfidf = TfidfVectorizer(min_df=2, ngram_range=(1,2))
        tfidf.fit(x_train)
        x_train_tf = tfidf.transform(x_train)
        x_dev_tf = tfidf.transform(x_dev)
```

0.1.2 2. Classifiers

We investigate the algorithms most used in today's machine learning community. Each classifier train on the training data and validate on the testing data.

```
In [4]: classifiers = {
        "K-Nearest Neighbors": KNeighborsClassifier(7),
        "Linear SVC": LinearSVC(),
        "Linear Classify with SGD": SGDClassifier(tol=1e-3),
        "Logistic Regression": LogisticRegression(),
        "MLP": MLPClassifier(),
        "Decision Tree": DecisionTreeClassifier(random_state=0),
        "Random Forest": RandomForestClassifier(random_state=0, n_estimators=100),
        "AdaBoost": AdaBoostClassifier(n_estimators=100),
        "Gaussian Naive Bayes": GaussianNB()
    }

    for nam, cla in classifiers.items():
        cla.fit(x_train_tf.toarray(), y_train)
        predicted = cla.predict(x_dev_tf.toarray())
        print("{0} - Accuracy: {1:.4f}".format(nam, metrics.accuracy_score(predicted, y_dev))
        print(metrics.classification_report(predicted, y_dev))
```

K-Nearest Neighbors - Accuracy: 0.6895

	precision	recall	f1-score	support
0	0.91	0.64	0.75	768
1	0.47	0.83	0.60	298
avg / total	0.78	0.69	0.71	1066

Linear SVC - Accuracy: 0.7880

	precision	recall	f1-score	support
0	0.78	0.80	0.79	527
1	0.80	0.78	0.79	539
avg / total	0.79	0.79	0.79	1066

Linear Classify with SGD - Accuracy: 0.7833

	precision	recall	f1-score	support
0	0.86	0.75	0.80	614
1	0.71	0.83	0.76	452
avg / total	0.79	0.78	0.78	1066

Logistic Regression - Accuracy: 0.7664

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.76	0.77	0.77	528
1	0.77	0.76	0.77	538
avg / total	0.77	0.77	0.77	1066

MLP - Accuracy: 0.7777

	precision	recall	f1-score	support
0	0.78	0.78	0.78	538
1	0.78	0.77	0.78	528
avg / total	0.78	0.78	0.78	1066

Decision Tree - Accuracy: 0.6107

	precision	recall	f1-score	support
0	0.62	0.61	0.62	540
1	0.61	0.61	0.61	526
avg / total	0.61	0.61	0.61	1066

Random Forest - Accuracy: 0.7186

	precision	recall	f1-score	support
0	0.72	0.72	0.72	541
1	0.71	0.72	0.71	525
avg / total	0.72	0.72	0.72	1066

AdaBoost - Accuracy: 0.6689

	precision	recall	f1-score	support
0	0.66	0.68	0.67	528
1	0.68	0.66	0.67	538
avg / total	0.67	0.67	0.67	1066

Gaussian Naive Bayes - Accuracy: 0.6895

	precision	recall	f1-score	support
0	0.73	0.68	0.70	578
1	0.65	0.70	0.67	488
avg / total	0.69	0.69	0.69	1066

0.1.3 3. Results

The only metrics used in the picture under is accuracy, and from the comparison of all classifiers, the svm with sgd optimization is the best one.

```
In [6]: import matplotlib.pyplot as plt
```

```
plt.rcParams()
fig, ax = plt.subplots()

models = classifiers.keys()
y_pos = np.arange(len(models))
performance = [0.6895, 0.7880, 0.7833, 0.7664, 0.7777, 0.6107, 0.7186, 0.6689, 0.6895]

ax.barh(y_pos, performance, align='center',color='green')
ax.set_yticks(y_pos)
ax.set_yticklabels(models)
ax.invert_yaxis() # labels read top-to-bottom
ax.set_xlabel('Accuracy')
ax.set_title('The classification preformance on different models.')

plt.show()
```

